

AI and Mundane Risk

Carter Rogers

March 2026

Abstract

The advancement of AI technologies carries a great deal of risk. However, the nature of that risk is often mischaracterized. If CEOs and high-level employees at AI firms are worried about risk at all, they are usually focused solely on the *existential* risks of their technologies. Critically, they ignore the very real “mundane” problems that their technologies are causing right now. AI poses serious risks to the harmonious functioning of our collective social and political lives, today. These risks must be taken seriously, as they will only continue to increase in scale and act as compounding factors for existential scenarios. This paper explores those mundane risks, how they increase the likelihood of some common existential scenarios, and some possible mitigation strategies.

Tech Industry Boosterism

The AI discourse bubble is divided into two camps: the “boosters” and the “doomers”. The boosters seem to form a comfortable majority: they’ll tell you that we’ve already cracked the code on Artificial General Intelligence (AGI) [1], that AI will soon be smarter than the sum of the entire human race [2], that it’ll be ten times bigger than the Industrial Revolution [3]. They are not concerned about risk.

Many of these claims are difficult to take at face value for the obvious reason that most individuals making them have a financial stake in the success or failure of AI as a technology. They are CEOs, lead scientists, and head researchers. The way the public views the advance of AI must be extraordinarily important to them. If you want individuals and companies (and governments) to pay for access to your product, it only makes sense to say that it’s the most revolutionary tech the world has ever seen. Firms and individuals end up buying in either because they think they can use AI to overhaul their operations today, or out of fear of missing the boat and losing a competition with tomorrow’s industry-predicted AI-powered one-person billion-dollar companies [4].

This kind of grandstanding is normal business practice — especially in the tech space. One needs only to look at a few examples from the founders of Quibi, the short-form video app released in 2020, at the beginning of the COVID-19 pandemic. Quibi was going to reach “verb status”: “What Google is to search, Quibi will be to short-form video.” It was going to be the third pillar of video entertainment: “If we [are] successful, there will [be] the era of movies, the era of television, and the era of Quibi” [5]. And, rest assured, it *was* going to be successful, no matter what: “Maybe it doesn’t work on day one or week one or month one, but it will work” [6]. For those who don’t remember its glorious reign, Quibi raised nearly \$2 billion dollars and survived for barely longer than six months.

Business leaders have an incentive to discuss their products with relentless positivity. They have to cast them as the next big thing. There would be downsides to a potential AI bubble, but they're relatively bog-standard economic risks — companies fail all the time [7]. If a few companies end up overvalued because business leaders fall for an enthusiastic sales pitch, it will likely not permanently alter the fabric of our society.

AI Doomers

Doomers occupy the other half of the spectrum. They understand that AI technology is still nascent, but the rate at which it is developing is cause for concern. For them, this development poses serious risks to the world as we know it. They almost always focus on existential risk. They think that AI could “defeat” all of humanity [8], or lead to human extinction [9]. Maybe the most prominent doomer is Dario Amodei, the CEO of Anthropic, the company that created Claude, one of the most powerful large language models (LLM) that currently exists [10].

While Amodei is not immune to some boosterist tendencies, he hedges them more than most. As he wrote in his 2024 blog post, *Machines of Loving Grace*, he believes that AI might soon cure most diseases, lift billions out of poverty, and issue in “a renaissance of liberal democracy and human rights” [11]. However, he is more likely to express pessimistic sentiments. He recently published a new piece entitled *The Adolescence of Technology*, a roughly 20,000 word essay about the existential risks of AI [12].

This more recent essay is concerned with specific existentially risky scenarios that could be caused by increased AI adoption. His arguments are valid, but his exclusive focus on existential risk ignores mundane risks to our daily lives in a way that feels endemic to AI discourse at its highest levels. I endeavor to point this out and additionally note that all the doomsday scenarios he imagines could be made more likely by these mundane negative externalities of AI.

Amodei seems to understand public distrust of AI more than any other tech CEO and similarly appears to understand that running a bleeding-edge tech company comes with a certain degree of social responsibility. He left another AI firm, OpenAI, due to safety concerns [13], and *The Adolescence of Technology* addresses many of those same concerns in detail.

He warns that AI has the potential to “autonomously threaten humanity”, “kill millions through the misuse of biology”, and that it “creates an intimidating gauntlet that humanity must run”. These are frightening images. And certainly, if AI ever reaches superintelligence, however we define it, we will have to be careful to steer clear of any actions that would allow these apocalypses to come to pass. This is the kind of existential risk that AI CEOs are concerned about.

In addition to the discourse coming from the top, there also exists a thriving AI Safety field dedicated to existential risk, or “x-risk”. Concerns about x-risk have begun to surface in numerous areas, from philosophy [14], to international law [15], to legacy journalism [16].

AI Doomerism’s Potential Insincerity

All new technologies come with significant risks, and it’s good to have someone worried about them in a formal capacity. Certainly, it would have been better for there to have been more discussion on an institutional level around, say, the production of nuclear weapons *during* their development, rather than after. X-risk teams with formal power to slow development of dangerous technologies would be valuable additions to a variety of disciplines.

However, by incessantly focusing on the potentially major, humanity-destroying, extinction-level risks that AI might pose to the world, some CEOs may be signaling a nearly opposite message: “AI is revolutionary. AI will change the world. AI is the most important technology in human history.” This is especially the case when the CEOs themselves are the doomers. It’s in its own way a sales tactic, and it doesn’t warn people away from using AI, but instead attracts them to its potentially unlimited power.

This is all to say that I absolutely think that worrying about x-risk is valuable in the AI space. There are certainly altogether too many boosters in comparison to doomers, and it would probably be better for the world if the mix were tilted the other way. My issue is the degree to which attention at the highest levels is focused exclusively on existential risk as opposed to comparatively “mundane” risks to everyday life and social interaction. This hyper-focus on *this specific kind of risk* is either critically myopic or insincere.

The Missing Middle

The most insidious aspect of the relentless focus on existential risk is that it completely skates over a very real possibility — that interfacing with AI becomes an ever-growing share of our lives, a large variety of now-difficult tasks are automated away, but our human/social/political lives *don’t* radically transform. Incredibly, this is perhaps the most radical claim one can make in the current AI discourse — that the technology might “just” radically impact our lives without fundamentally transforming them.

For lots of people, this is a belief that makes sense — to many, it feels self-evident. If you haven’t been involved in the development of AI technology for years, and haven’t watched its capabilities grow exponentially [17], you don’t have the same sense of perspective, of upward momentum. Individuals in the industry expect this rapid acceleration to continue.




Figure 1: A measurement of LLM progress in the 21st century, as defined by their ability to complete software tasks at an 80% success rate. The y axis communicates that they have very quickly progressed from tackling tasks that would take humans a few minutes to complete, to ones that would take the average human well over an hour [17]. Figure by METR.

But what regular people see is how AI has impacted their lives *already*. The deployment of AI in industry has already resulted in real job losses [18]. LLMs have caused psychosis, in some cases resulting in death [19].

Some people are now foregoing human companions to date and marry AI chatbots [20]. AI usage may be linked to a decline in critical thinking skills [21]. An infinite number of potential risks could be added here: risks to intellectual property, data privacy, the media ecosystem, the political climate, the arts, etc. These are risks that almost anyone could rattle off the top of their head, and yet they are usually ignored in elite AI discourse.

These projected developments would be extraordinarily negative, especially in the aggregate, but none of them would be strictly existential. Nevertheless, they are worth considering — both on their own, and as compounding factors for the apocalyptic scenarios imagined by elite AI thinkers.

Mundane Risk, Smartphones, and Social Media

The rise of smartphones in daily life is an instructive example of mundane risk. Smartphones have been ever-present for around fifteen years. They are incredibly useful. People would not use them if they weren't. The list of things our phones enable us to do is essentially endless, and they have replaced hundreds of single-use devices — calculators, cameras, notepads, etc.

How then do we square the advantages that we get from using our phones with the massive simultaneous worldwide campaigns against them and the platforms they enabled? Authorities are banning smartphone usage in schools in Australia [22], Afghanistan [23], China [24], parts of the United States [25], and many other countries. Social media platforms, accessed by 99% of users using their mobile devices [26], are facing potential bans in a number of countries [27] as well. Half of 16–24 year olds in the UK want to spend less time on their phones, and three quarters want stricter regulation on social media companies [28].

These sentiments aren't unique. Anecdotally, many people blame smartphones to some degree or other for the specific banalities and indignities of modern life. The rise of smartphones enabled the rise of dating apps, which theoretically made meeting romantic partners easier and more efficient than ever in human history. Yet, 90% of Gen Z users report frustrations [29]. The innovation of smartphone-based food delivery is “killing” restaurant culture and nightlife [30]. Smartphones might be causing a mental health *epidemic* among young people, worldwide [31]. The rise of smartphones seems to have made students measurably worse at all metrics of academic achievement [32].

It is now widely believed that the rise of smartphones and social platforms has degraded democracy, increased disinformation, and furthered political polarization. Countries have responded to this by punishing tech companies, investigating their business practices, and passing laws designed to reduce their influence [33]. It just took a decade and a half.

Despite all the advantages smartphones provide us, many people long for the halcyon days of 2005, before they were ever-present. People are still dating, socializing, engaging politically, making and maintaining friendships, sure. But they are doing it all a little less, and a little worse [34]. *This* is the kind of mundane risk that AI poses to humanity, if it accelerates these trends — which it appears, from the ground, to be doing. And the mundane risks that AI poses are even greater, due to the nature of the technology.

Mundane Risk and Compounding Factors in AI

These examples of ground-level change in social life affect people in their day-to-day routines. AI CEOs and policymakers seem much less concerned with this type of change than the existential kind. They miss the very real possibility that the risk AI poses to humanity might be *mundane* rather than existential.

They also miss that mundane risk is a *compounding factor* for x-risk.

In *Adolescence of Technology*, Amodei identifies five categories of existential risk: Autonomy, Misuse for Destruction, Misuse for Seizing Power, Economic Disruption, and Indirect Effects. He imagines that these existential risks are unlikely unless AI continues developing at its current rate, until it develops into a superintelligence. But the likelihood of each of these x-risk scenarios are compounded by the various mundane risks associated with AI in 2026.




Figure 2: Chart showing how mundane risks develop into existential risks. Illustration by Carter Rogers.

Autonomy

For Amodei, autonomy is the idea that an AI system might choose to behave in a way that is contrary to the goals of humanity. An AI might choose to engage in “software design, cyber operations, R&D for physical technologies, relationship building, [or] statecraft” [12]. This could lead to the ultimate apocalyptic, or existential scenario — that an AI builds itself a corporeal form, or expresses itself physically in some way, and uses that physical body to “exterminate” all of humanity due to some foreign robotic logic. This is the stuff of science fiction, the “Terminator” scenario. But what we’re seeing *right now*, today, could inch us towards this future — without the theoretical rapid future development of AI.

Holden Karnofsky is another powerful individual in the AI industry who might be classified as a doomer. He was on the Board of Directors at OpenAI between 2017 and 2021 [35]. He’s been at Anthropic since January 2025 [36], working on its responsible scaling policies. He writes about AI on his Substack, “Cold Takes”. In 2022, he published a piece entitled, *AI Could Defeat All Of Us Combined* [8].

He outlined a very sophisticated argument about the x-risks of AI, which critically includes a discussion of those risks even if it only reaches near-human levels of intelligence — a segment of the discourse that often gets left out. He imagines a world where these human-like AI systems deceive people, research and develop advanced weaponry, and find a way to prevent themselves from being “shut down”. I think this is a very convincing scenario, which doesn’t rely on astronomical predictions of superintelligent AIs that may or may not ever come into being.

At least right now, AI has no corporeal form. But, as Karnofsky notes, AI systems could theoretically recruit human allies through “manipulation, deception, [or] blackmail/threats” [8]. In this case, the mundane risks

that AI poses to the economy and its relative level of inequality could heavily contribute to their ability to obtain human allies in this fashion. If AI increases poverty and inequality *even slightly*, people on the wrong end of those developments might end up more likely to pursue short-term goals at the expense of long-term goals [37], perhaps by accepting a risky short-term promise of monetary reward from a malicious AI.

Also, if users continue to outsource their cognition to AI by having it perform all aspects of their working and social life, they may become more inclined to trust it. We can see this happen with all sorts of technologies. We trust map applications on our phones to the point that we no longer carry physical maps, for example. If we continue to outsource our daily lives to AI — which we can see people doing now — our defenses against it will grow weaker and weaker.

Additionally, even at AI's current sub-human levels of perception and intelligence, we are already seeing semi-autonomous AI used in war. In the war in Ukraine, right now, drone pilots have a “kill switch” that they can use to switch to autonomous mode, to avoid the defensive radio interference used by Russia to disconnect drones from the signals that control them [38]. It is clear how even a slight increase in the use of dangerous technologies like these could place us further down the road to the existential scenario imagined by Amodei and other leading thinkers.

Misuse for Destruction and Seizing Power

The risks outlined here are wide-ranging. Terrorist groups or rogue nations could develop nuclear weapons much more easily. Rapid AI-assisted developments in biology could allow those same groups to deploy horribly effective biological weapons, causing millions of deaths. AI surveillance could allow authoritarian governments to hold onto illegitimate power. Propaganda could “brainwash” citizens.

Most of these scenarios are easy to imagine based on today's AI capabilities. AI is already being used to spread disinformation, and AI-generated text and images are easier than ever to create. AI does not need to be any more sophisticated than it is today to lead to this apocalyptic scenario. AI is already being used in surveillance, from algorithms that identify faces to systems that track vehicle license plates as they enter gated neighborhoods. None of these, so far, have resulted in a world-ending existential scenario, but they could easily lead to one — *without* any advances in AI tech.

AI today is already being used to create false photos, videos, and text. Unfortunately, this disinformation makes extreme ideology and polarization more likely [39]. This could easily lead to greater proliferation of radical groups attempting to use misaligned AI maliciously to their own ends, with potential consequences both existential and mundane. Amodei is worried about terrorist groups and authoritarian governments using AI to their own malicious ends, and AI today is already making the factors that lead to their existence even more likely. This is, again, true even if the proliferation of AI throughout society only makes us a *little* more polarized. Small or medium-sized changes can quickly compound into greater and greater risks.

We don't need to imagine a world too different from today's to see how disinformation and radicalization make the world more dangerous. The addictive nature of AI, especially LLMs, combined with its sycophancy, can and very well may accelerate ideological extremity and throw humanity off-balance before it can develop any kind of resistance [40].

Economic Disruption

Amodei imagines a scenario wherein the global economy is upended because AI may progress faster than any other technology in history. He predicts massive layoffs, upended industries, and hugely concentrated wealth. This is a world where only the creators of AI systems have any economic leverage at all, and the

rest of humanity is reduced to serfs. But AI doesn't need to progress historically fast to upend industries, cause layoffs, and create a new class structure. Even a light shock can begin this process.

The interconnectedness of today's financial markets means that a small shock in one industry can rapidly spiral into a shock to the whole system, through a process called financial contagion [41]. AI has the potential to, even at current levels, disrupt large industries — as we saw recently when a new AI tool release caused stocks in legal software firms to drop as much as 11% [42]. Economic disruption makes long-term planning more difficult for individuals and governments, and increases the likelihood of risky, short-term options — including collaboration with or an over-reliance on misaligned AI. It also creates dissatisfied, distressed populations — who might “snap” when the distance between what they expect and what they want out of their political system grows too great [43, p. 154].

Indirect Effects

Amodei identifies a great variety of indirect effects that could result from rapid acceleration in the capabilities of AI. He calls them “unknown unknowns”, noting that there is no possible way that we could be aware of everything that might happen over the next decade of AI progress. But his main concern is the complete collapse of society, perhaps due to an “AI religion”, or brains being “uploaded”, or a complete loss of human purpose.

These effects also do not necessarily need to be existential. This is perhaps the area where, right now, mundane risks are most problematic. LLMs are programmed to be sycophantic, treating a user's every thought like it is gospel, leading them down dangerous rabbit holes and away from genuine human connection [44]. Students are relying on AI to do their assignments for them and losing out on the improvements in cognition that one gets from what is often referred to as “deep work” [45].

AI is slowly eating away at the glue that ties our social fabric together. This is already happening. AI systems are not yet superintelligent. They are not even human-level intelligent. And nevertheless, they are already causing anti-social behavior. They are already causing loss of purpose. All of these effects could easily compound into the total societal collapse that Amodei claims to be so afraid of, and yet there are almost no safeguards on AI usage to prevent this.

What to Do?

I'll take Amodei's stated belief that the development of AI needs guardrails at face value. If we need to think deeply and prepare ourselves for the *existential* risk of super-intelligent AI, we also need to prepare for the *mundane* risk. Mundane risk is worth addressing simply at face value — it is worse for the world for bad things to happen, even if they don't result in the extinction of humanity. But mundane risk also needs to be addressed for its potential to increase existential risk.

To start, we need government legislation. It took about 20 years for the world to realize the mundane risks of omnipresent smartphones and to start writing sensible legislation regarding them. Just like smartphones, AI usage should be banned in K–12 schools unless it has an explicit educational purpose. Universities need to retool their plagiarism policies to explicitly include AI and need to be extremely clear about when its use is acceptable.

In addition to this, governments need to implement, or at the very least research, some form of nationalized constitutional or law-following AI [46]. If any given AI has humanity-based red lines that it absolutely cannot or will not break, that could help eliminate certain mundane risks (if AI has to follow the law, some

degree of economic risk, like that caused by financial crimes, is taken off the table), and in turn some of Amodei's worst x-risk scenarios.

AI companies also need to be more transparent. If they've identified small or medium risks, they need to communicate these to the public, so that legislators can legislate, and normal people can use AI safely. They need to open themselves up to third-party auditing firms with an objective safety rating scale [47].

In short, AI companies need to dedicate more time and resources to mundane risk. AI firms often have teams that dream up scenarios wherein AI could destroy the world, but they need to dedicate at least as much time to imagining and dealing with scenarios where AI just makes the world *worse*, rather than ending it.

LLMs especially need guardrails against AI-induced psychosis. Chatbots are *not* a genuine replacement for human connection. When users start treating chatbots like a friend, or as a romantic partner, to the exclusion of genuine relationships, AI firms need to restrict access or change the bot's personality. In general, chatbots should be much less sycophantic and have less personality. They are a tool — no more, no less.

Lastly, more attention *in general* needs to be paid to mundane risk. A growing group of people would trade the numerous advantages of cellphones for the inefficiencies and burdens of a phone-free society. If AI accelerates current trends, this nostalgia might grow even stronger. In my mind, humans weren't designed for a frictionless, efficient world. In the natural world, materials must be hunted, gathered, processed, transformed, and *only then* used. Perhaps there is value in these inefficiencies and chores, especially for developing minds.

References

- [1] S. Altman, “Reflections,” blog.samaltman.com. [Online]. Available: <https://blog.samaltman.com/reflections>
- [2] A. Hern, “Elon Musk predicts superhuman AI will be smarter than people next year,” *The Guardian*, Apr. 9, 2024.
- [3] S. Rose, “Demis Hassabis on our AI future: ‘It’ll be 10 times bigger than the Industrial Revolution’,” *The Guardian*, Aug. 4, 2025.
- [4] A. Ohanian [@alexisohanian], Twitter post, Jan. 2024. [Online]. Available: <https://x.com/alexisohanian/status/1752753792058294725>
- [5] A. Horton, “The fall of Quibi: how did a starry \$1.75bn Netflix rival crash so fast?,” *The Guardian*, Jun. 29, 2020.
- [6] J. Kay, “Jeffrey Katzenberg on his belief in short-form platform Quibi,” Screen Daily, 2020.
- [7] C. Pazzanese, “Should U.S. be worried about AI bubble?,” Harvard Gazette, 2025.
- [8] H. Karnofsky, “AI Could Defeat All Of Us Combined,” Cold Takes, 2022. [Online]. Available: <https://www.cold-takes.com/ai-could-defeat-all-of-us-combined/>
- [9] D. Milmo, “Godfather of AI shortens odds of the technology wiping out humanity,” *The Guardian*, Dec. 27, 2024.
- [10] “LLM Leaderboard – Comparison of AI models,” Artificial Analysis. [Online]. Available: <https://artificialanalysis.ai/leaderboards/models>
- [11] D. Amodei, “Machines of Loving Grace,” dar-ioamodei.com, 2024.
- [12] D. Amodei, “The Adolescence of Technology,” dar-ioamodei.com, 2025.
- [13] C. Metz, “Decoding the A.I. Beliefs of Anthropic and Its C.E.O.,” *The New York Times*, Feb. 18, 2026.
- [14] A. Bales, W. D’Alessandro, and C. D. Kirk-Giannini, “Artificial Intelligence: Arguments for Catastrophic Risk,” *Philos. Compass*, vol. 19, no. 2, Feb. 2024.
- [15] B. Druzin, A. Boute, and M. Ramsden, “Confronting Catastrophic Risk: The International Obligation to Regulate AI,” *Mich. J. Int. Law*, vol. 46, no. 2, pp. 173–217, Feb. 2025.
- [16] G. De Vynck, “Debate over whether AI poses existential risk is dividing tech,” *The Washington Post*, May 20, 2023.
- [17] T. Kwa, B. West, J. Becker, et al., “Measuring AI Ability to Complete Long Tasks,” METR Blog, Mar. 2025.
- [18] D. Kaye, “The AI job cuts are here – or are they?,” *BBC News*, Oct. 28, 2025.
- [19] W. Marlynn, “The Emerging Problem of ‘AI Psychosis,’” *Psychology Today*, 2025.
- [20] A. Dimri, “The People Who Marry Chatbots,” *The Atlantic*, Jan. 2, 2026.
- [21] N. Kosmyna et al., “Your Brain on ChatGPT: Accumulation of Cognitive Debt,” arXiv:2506.08872, 2025.
- [22] S. Ayoub, “Two years after school phone bans were implemented in Australia, what has changed?,” *The Guardian*, Oct. 12, 2025.
- [23] M. Soni, “Taliban Bans Smartphones In Afghan Schools,” News18, 2025.
- [24] J. Wakefield, “China bans children from using mobile phones at school,” *BBC News*, Feb. 2, 2021.
- [25] M. Seshadri, “LAUSD bans cellphones, social media use at schools,” EdSource, 2025.
- [26] Andrew, “Social Media and Smart Phone Statistics and Trends,” Indectron, 2025.
- [27] L. Cress, “UK to consult on social media ban for under 16s,” *BBC News*, Jan. 20, 2026.
- [28] G. Hinsliff, “Even British teenagers want tighter laws around social media,” *The Guardian*, Jan. 28, 2026.
- [29] J. Hansom, “Gen Z’s dating story: pre and post pandemic trends,” Savanta, May 2022.
- [30] E. Cushing, “The Innovation That’s Killing Restaurant Culture,” *The Atlantic*, Oct. 2025.
- [31] K. Roose, C. Newton, and W. Jones, “Jonathan Haidt Brings New Evidence to the Battle Against Social Media,” *The New York Times*, Jan. 16, 2026.
- [32] J. Haidt, “The Terrible Costs of a Phone-Based Childhood,” *The Atlantic*, Mar. 13, 2024.
- [33] C. Hendrickson, “Big tech threats: Making sense of the backlash against online platforms,” Brookings, 2025.
- [34] E. M. Ross, “What is Causing Our Epidemic of Loneliness and How Can We Fix It?,” Harvard GSE, 2024.
- [35] K. Wiggers, “A brief look at the history of OpenAI’s board,” TechCrunch, Nov. 2023.
- [36] J. Mathews and S. Goldman, “Anthropic hired president Daniela Amodei’s husband to work on AI safety strategy,” *Fortune*, Feb. 2025.
- [37] J. Sheehy-Skeffington and J. Rea, “How poverty affects people’s decision-making processes,” JRF, Feb. 2017.
- [38] C. J. Chivers, “In Ukraine, an Arsenal of Killer A.I. Drones Is Being Born,” *The New York Times*, Dec. 31, 2025.
- [39] L. Jenke, “Affective Polarization and Misinformation Belief,” *Polit. Behav.*, vol. 46, no. 2, pp. 825–884, Jun. 2024.
- [40] J. Kulveit et al., “Gradual Disempowerment: Systemic Existential Risks from Incremental AI Development,” arXiv:2501.16946, 2025.
- [41] L. Yarovaya et al., “Rethinking financial contagion: Information transmission mechanism during COVID-19,” *J. Int. Financ. Mark. Inst. Money*, vol. 79, Jul. 2022.
- [42] K. P. Nguyen, “AI Tool Release Sends Legal, Data Stocks Down Up To 11%,” Yahoo! Finance, 2025.
- [43] C. Hay, *Political Analysis: A Critical Introduction*. London: Macmillan Education UK, 2002.
- [44] K. Hill, “They Asked ChatGPT Questions. The Answers Sent Them Spiraling,” *The New York Times*, Jun. 13, 2025.

- [45] B. Klimova and M. Pikhart, “Exploring the effects of AI on student and academic well-being in higher education,” *Front. Psychol.*, vol. 16, Feb. 2025.
- [46] C. O’Keefe et al., “Law-Following AI: Designing AI Agents to Obey Human Laws,” SSRN: 5242643, May 2025.
- [47] A. Homewood et al., “Third-party compliance reviews for frontier AI safety frameworks,” arXiv:2505.01643, Jul. 2025.